

Introduction to Reconfigurable Computing

Introduction to the development kits and design software

LECTURE 1

IOULIIA SKLIAROVA

Reconfigurable Computing

A computer architecture that facilitates faster and more complex computing by using reprogrammable integrated circuits to process data that solves problems and directs specified functions.

Offers the higher speed and efficiency of computer hardware to perform serial and parallel tasks combined with software's flexible capacity to allow changes to a circuit's purpose in the field to adjust for changing circumstances.

Core advantages:

- More functionality from simpler and smaller hardware designs
- Cost savings on low volume products and those whose useful life is extended by updating
- Shorter/faster development time-to-market

Original Idea

The earliest work on reconfigurable computer architecture was done at the University of California at Los Angeles – in 1960's Gerald Estrin proposed a standard processor extended with an array of “variable” hardware.

The main goal was to permit computations which are beyond the capabilities of present systems by providing an inventory of high speed substructures and rules for interconnecting them such that the entire system may be temporarily distorted into a problem oriented special purpose computer.

At that time, digital technology was not ready for such a revolutionary change.

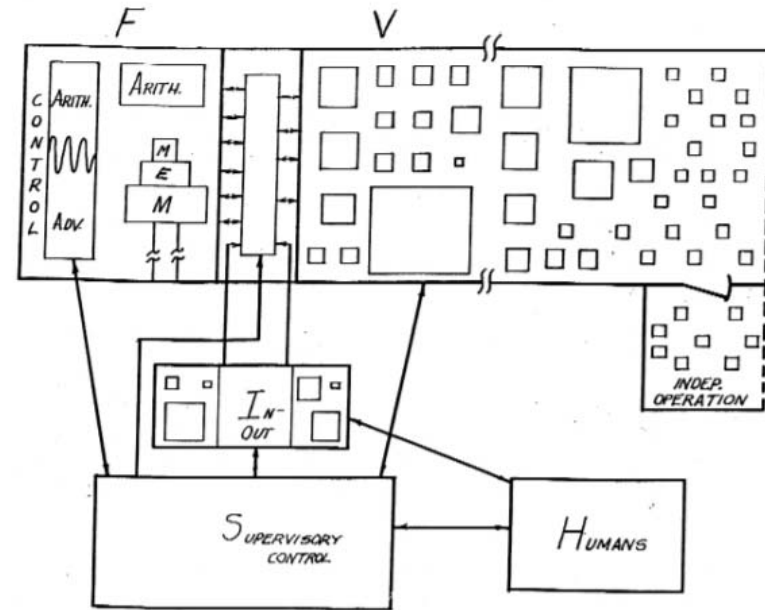


Figure 1. This diagram shows the relations between the fixed machine, the variable structure part, the input-output, the supervisory control, and the humans. (Courtesy of the 1960 Western Joint Computer Conference.)

Other Computing Models

Traditionally, computing was classified into General-Purpose Computing performed by a **General-Purpose Processor (GPP)** and Application-Specific Computing performed by an **Application-Specific Integrated Circuit (ASIC)**.

Reconfigurable computing acts as a trade-off between the two extreme characteristics of GPP and ASIC, combining the advantages of both.

When compared to GPP:

- provides ability to make substantial changes to the datapath itself in addition to the control flow
- can have better performance with respect to a software implementation in GPP but paying this in terms of time to implement

When compared to ASIC:

- possibility to adapt the hardware <during runtime> by "loading" a new circuit on the reconfigurable fabric
- can be used to design a system without requiring the same design time and complexity compared to a full custom solution but being beaten in terms of performance

Overview

Advantages / Features

- Flexible hardware/software co-design
- Customization (parallelism exploration, specialized control and datapaths)
- Fast prototyping
- Faster design cycles, lower risk, cheaper for small to medium size markets compared with ASICs
- Dynamic system upgrades (field updates for both hardware and software)

Target applications

- Application-specific hardware accelerators/coprocessors/processors
- Telecommunications, networks, image/audio processing, cryptography, space, automotive, medical, etc.

Overview

Limitations

- Overhead due to reconfiguration time and/or additional hardware circuit delays


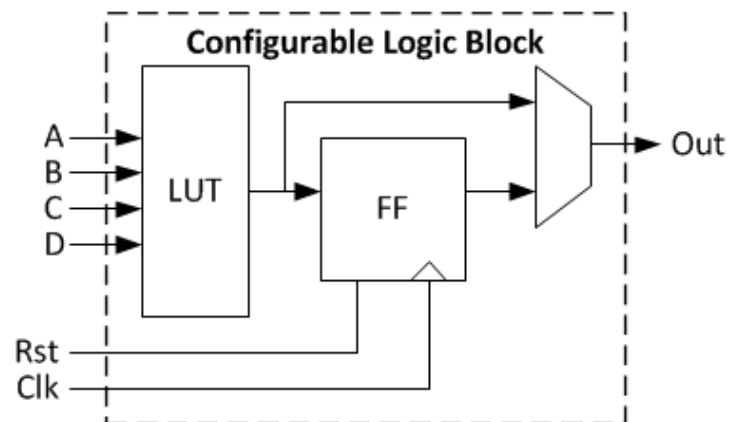

Supporting technologies

- High capacity FPGAs/Programmable SoC
- Processor customization frameworks
 - Fixed processor cores + extension coprocessors / flexible processors

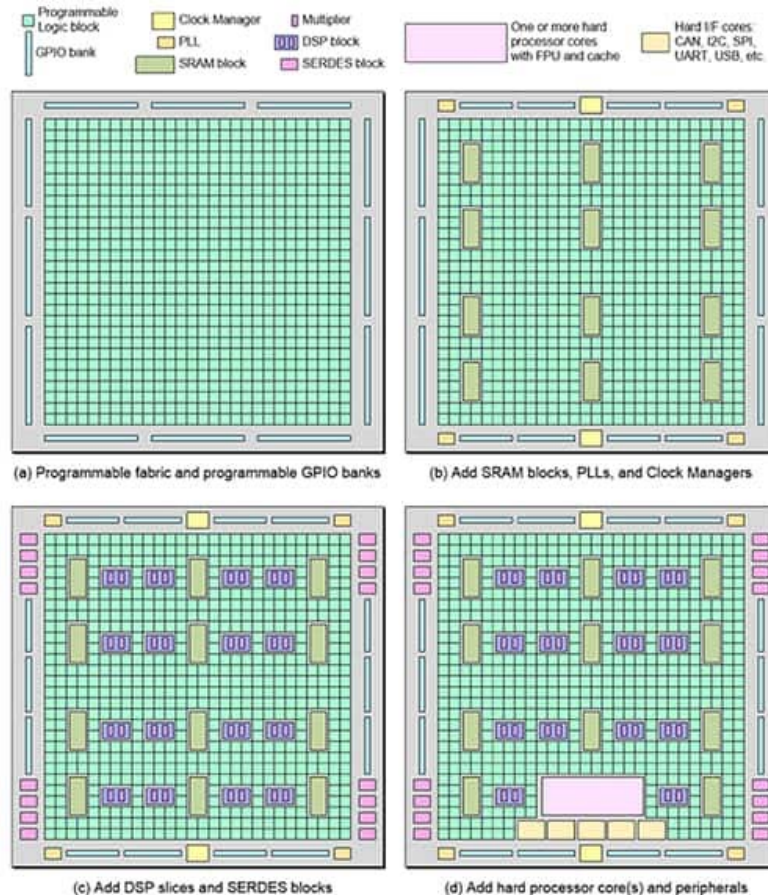
Reconfiguration approaches

- Static (at manufacturing, deployment or boot time)
- Dynamic
 - Total (at run-time between execution phases)
 - Partial (at runtime replacing some components, while others remain operational)

FPGA – Field-Programmable Gate Array



Modern FPGA



The simplest FPGAs contain only programmable fabric and configurable general purpose IO (GPIO) (a); different architectures augment this fundamental fabric with memory blocks, PLLs, and clock managers (b); DSP blocks and SERDES interfaces (c); and hard processor cores and peripherals (d). (Image source: Max Maxfield)

6-input LUTs configurable as distributed memory

Embedded memory blocks - block RAM with built-in FIFO logic for on-chip data buffering

High-speed serial connectivity with built-in multi-gigabit transceivers

User configurable analog interfaces, incorporating analog-to-digital converters

DSP slices with dedicated multipliers for high-performance filtering

Clock management tiles providing clock frequency synthesis, deskew, and jitter filtering functionality

Soft and hard processors

- ⇒ FPGAs with hard processor cores are referred to as programmable systems-on-chip (PSoC)

Xilinx Programmable Devices

The performance and capabilities of the programmable device offerings from Xilinx span from modest to extremely high:

- traditional FPGAs
- PSoCs (FPGA programmable fabric with a single hard core processor)
- MPSoCs (FPGA programmable fabric with a multiple hard core processors)
- RFSoc (MPSoCs with RF capability)
- ACAPs (Adaptive Compute Acceleration Platforms)


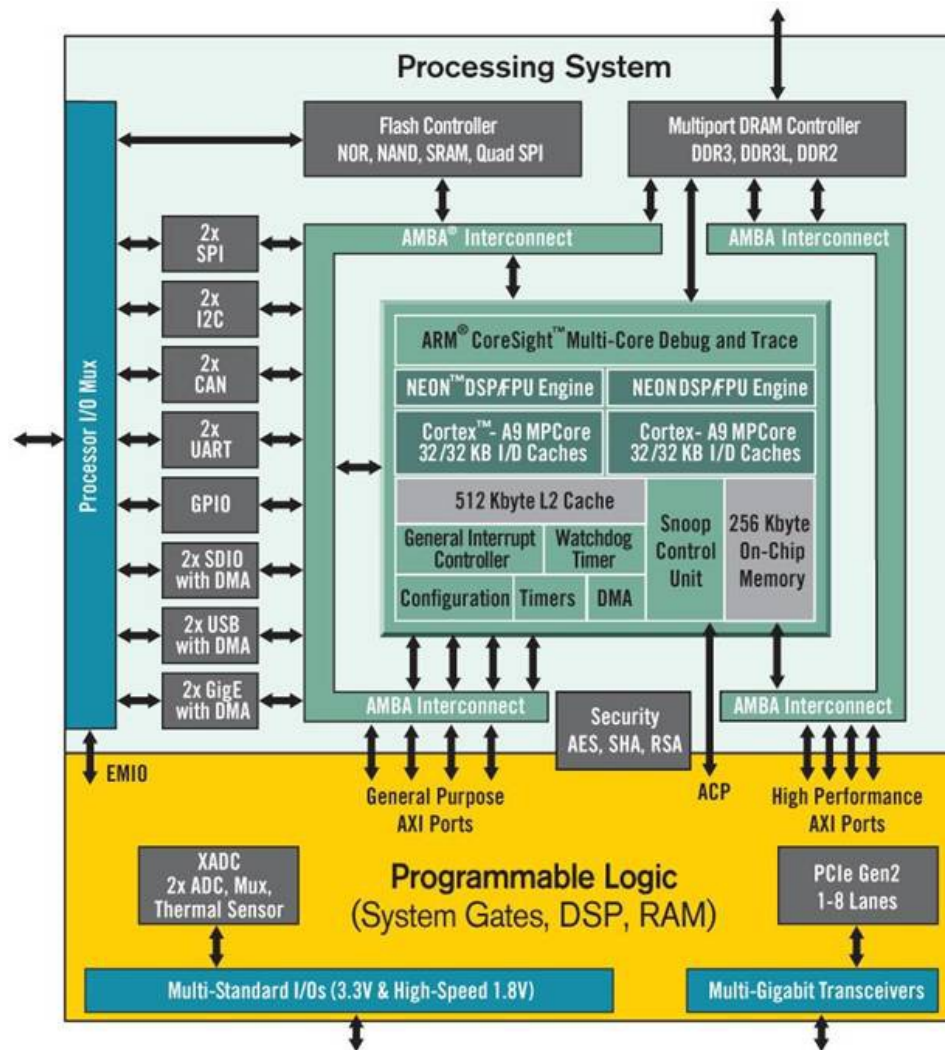


Figure 2: Over time, the Xilinx architectural portfolio has evolved from simple FPGAs containing only programmable fabric, to SoC devices in which the programmable fabric is augmented with a hard core processor, to MPSoCs with multiple processors, to RFSoc with RF capabilities, to the latest generation of ACAPs, which are targeted toward applications like AI. (Image source: Max Maxfield)

PSoC example – Zynq-7000



Xilinx FPGA




Figure 3: Xilinx FPGA offerings provide a comprehensive multi-node portfolio to address requirements across a wide set of applications. (Image source: Max Maxfield)

Table 4: Artix-7 FPGA Feature Summary by Device

Device	Logic Cells	Configurable Logic Blocks (CLBs)		DSP48E1 Slices ⁽²⁾	Block RAM Blocks ⁽³⁾			CMTs ⁽⁴⁾	PCIe ⁽⁵⁾	GTPs	XADC Blocks	Total I/O Banks ⁽⁶⁾	Max User I/O ⁽⁷⁾
		Slices ⁽¹⁾	Max Distributed RAM (Kb)		18 Kb	36 Kb	Max (Kb)						
XC7A12T	12,800	2,000	171	40	40	20	720	3	1	2	1	3	150
XC7A15T	16,640	2,600	200	45	50	25	900	5	1	4	1	5	250
XC7A25T	23,360	3,650	313	80	90	45	1,620	3	1	4	1	3	150
XC7A35T	33,280	5,200	400	90	100	50	1,800	5	1	4	1	5	250
XC7A50T	52,160	8,150	600	120	150	75	2,700	5	1	4	1	5	250
XC7A75T	75,520	11,800	892	180	210	105	3,780	6	1	8	1	6	300
XC7A100T	101,440	15,850	1,188	240	270	135	4,860	6	1	8	1	6	300
XC7A200T	215,360	33,650	2,888	740	730	365	13,140	10	1	16	1	10	500

Notes:

- Each 7 series FPGA slice contains four LUTs and eight flip-flops; only some slices can use their LUTs as distributed RAM or SRLs.
- Each DSP slice contains a pre-adder, a 25 x 18 multiplier, an adder, and an accumulator.
- Block RAMs are fundamentally 36 Kb in size; each block can also be used as two independent 18 Kb blocks.
- Each CMT contains one MMCM and one PLL.
- Artix-7 FPGA Interface Blocks for PCI Express support up to x4 Gen 2.
- Does not include configuration Bank 0.
- This number does not include GTP transceivers.

Artix-7 CLB (Configurable Logic Block)

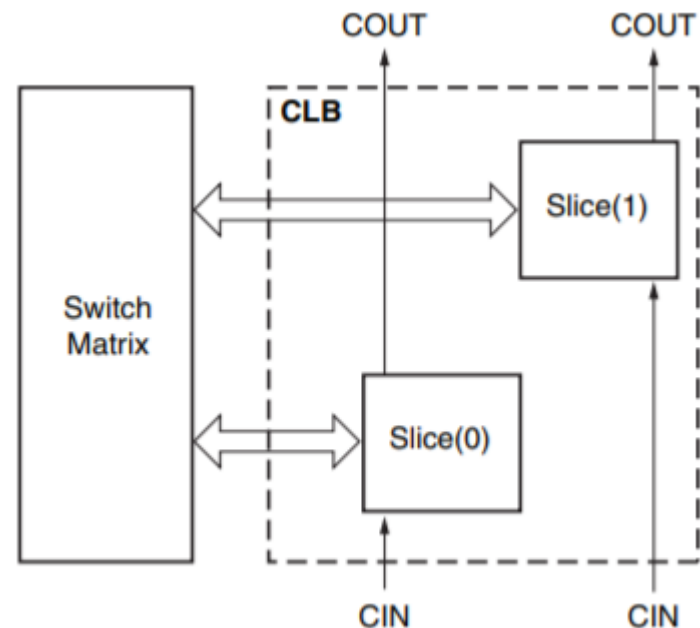
CLBs are the main logic resources for implementing sequential as well as combinatorial circuits.

Each CLB element is connected to a switch matrix for access to the general routing matrix.

A CLB element contains a pair of slices.


Every slice contains:

- Four logic-function generators (or look-up tables)
- Eight storage elements
- Wide-function multiplexers
- Carry logic




Nexys-4 Development Board

- 16 user switches
- 16 user LEDs
- Two 4-digit 7-segment displays
- USB-UART Bridge
- Two tri-color LEDs
- Micro SD card connector
- 12-bit VGA output
- PWM audio output
- PDM microphone
- 3-axis accelerometer
- Temperature sensor
- 10/100 Ethernet PHY
- 16MB CellularRAM
- Serial Flash
- Four Pmod ports
- Pmod for XADC signals
- Digilent USB-JTAG port for FPGA programming and communication
- USB HID Host for mice, keyboards and memory sticks



FPGA: xc7a100Tcsg324-1

Nexys-4 Development Board



Development Tools

Xilinx Vitis

- For portability of projects between different computers, it is recommended to install **version 2020.2**
- Installation instructions are available at the course website
- Requires at least **70 GB** of disk space
- License: Vivado Design Suite 30-day evaluation license -> floating license available within the UA campus or with active VPN





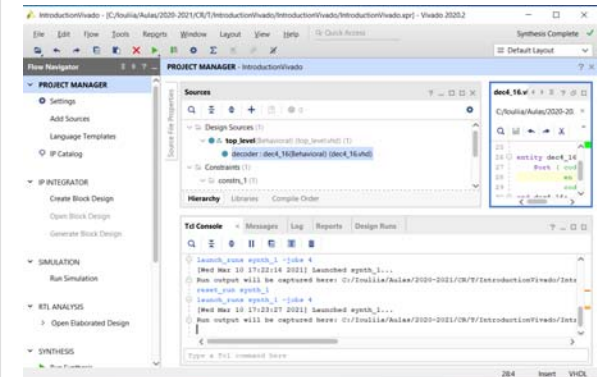
Figure 6: A high-level view of the Xilinx Vivado and Vitis design tool stack reflects how users can work with the tools at the most appropriate levels of abstraction. Hardware designers work with Vivado, software developers work with Vitis, and AI and data scientists work with Vitis AI. (Image source: Max Maxfield)

FPGA Design Flow




Design entry based on:

- Hardware description languages
- State diagrams
- Schematic capture
- C/C++ and other high level languages



Logic Synthesis




Results in a netlist (i.e. hardware components and their interconnections) that implement the modeled behavior and structure

Outputs

- Netlist
- Circuit performance and resource usage estimations

Implementation (Map, Place and Route)



Maps the netlist in FPGA primitives


Places the primitives in specific FPGA locations

Routes the interconnections between primitives

Outputs

- FPGA configuration file
- Reports with FPGA resource usage, circuit delays, power consumption, ...

Device (FPGA) Programming



Transfers the configuration file to the FPGA

- Done with a software tool and a programming cable
- Most FPGAs are based on SRAM (volatile configuration)
- Some devices are based on non volatile FLASH memory

System-on-Chip (SoC)

A **SoC** is a single integrated circuit integrating at least components like a processor and I/O peripherals to perform general user interface and housekeeping tasks and special hardware accelerators to handle computation-intensive operations.

A **SoC** integrates a processor, memory modules, I/O peripherals, and custom hardware accelerators into a single chip.

We will construct several SoCs containing a Xilinx MicroBlaze soft-core processor, memory-mapped I/O subsystem that can incorporate custom I/O cores,

Hardware-Software Co-Design

The FPGA technology allows us to tailor the processor, select only the needed I/O peripherals, create a custom I/O interface, and develop specialized hardware accelerators for computation-intensive tasks.

Both the hardware and software can be customized to match specific needs.

The methodology of exploiting the trade-offs between hardware and software and developing and integrating them concurrently is referred to as **hardware-software co-design**.

Development Flow of SoC

- Partition the tasks to software routines and hardware accelerators
- Design user custom cores if needed
- **Develop the hardware**
- Develop the software
- Implement the hardware and software and perform testing.

Final Remarks

At the end of this lecture you should be able to:

- Describe the main features of reconfigurable computing
- Outline a typical FPGA architecture
- Know difference between FPGA and PSoC
- Identify the main steps of FPGA design flow
- Define software-hardware co-design

To do:

- Install Xilinx Vitis 2020.2